

Mental Wealth Academy


Community Intelligence Network (CIN)

*A Decentralized Framework for Community-Owned Mental Wellness, Financial Literacy,
and Behavioral Science Research*

James Q. Marsh, Psych B.S.

research@mentalwealthacademy.net


Mentalwealthacademy.net | github.com/Mental-Wealth-Academy



Abstract

Framework Overview

Mental Wealth Academy (MWA) Community Intelligence Network (CIN) is a theoretical framework addressing the structural failures in behavioral health data: community-generated knowledge about mental wellness, financial behavior, and psychological resilience is routinely extracted by academic institutions and commercial platforms without returning value to the people who produced it. CIN proposes a three-layer architecture — (1) decentralized member-owned data stores, (2) a structured federated learning protocol coordinated on a public blockchain, and (3) a trustless cryptographic rewards mechanism — that enables MWA's member community to contribute to real research while retaining ownership of their data and receiving direct compensation for their participation. Blue, MWA's AI guide, functions as the human-facing interface to this infrastructure: facilitating curriculum delivery, structured assessments, and behavioral surveys while ensuring that the research pipeline serves the community rather than extracting from it. By eliminating institutional gatekeepers and realigning incentive structures, CIN positions MWA as a platform where members are simultaneously learners, researchers, and co-owners of the intelligence they collectively produce.




1. Introduction

The internet promised democratization. Instead, two decades of platform capitalism have produced the opposite: a world in which individuals generate enormous quantities of behavioral, psychological, and financial data, while the economic and scientific value of that data flows almost entirely to centralized institutions. Nowhere is this more apparent than in behavioral health — a domain where the lived experience of millions of people dealing with financial stress, trauma, anxiety, and systemic disadvantage has been transformed into academic publications and commercial products, none of which return material benefit to the communities studied.

Mental Wealth Academy (MWA) was built in rejection of this model. Founded as a mobile-first learning and research platform, MWA provides structured curriculum in mental wellness, financial literacy, and behavioral science to a community of members who are often underserved by traditional mental health and financial education systems. Members work through progressive curriculum tracks, complete validated behavioral assessments, participate in community research surveys, and interact with Blue — MWA's AI guide — in ways that generate real, analyzable data about psychological and financial wellbeing.

The core problem is this: that data, if handled through conventional research pipelines, would leave the community entirely. Universities would publish it. Platforms would sell it. The members who produced it would receive nothing. CIN is designed to change that.

Community Intelligence Network (CIN) is a theoretical framework describing how MWA can implement decentralized federated learning, blockchain-coordinated rewards, and member-owned data stores to make research participation a genuine economic and scientific contribution — one that funds the community instead of extracting from it. This paper describes the problem in depth, outlines the CIN architecture, details the protocol implementation, and situates MWA's approach within the broader landscape of decentralized AI research.



2. Problem Statement

MWA members navigate overlapping and compounding challenges. Financial stress and mental health are deeply intertwined — research consistently shows that financial precarity is both a cause and a consequence of psychological distress [1]. Yet the systems designed to study and address these dynamics replicate the very power asymmetries that produce them. CIN identifies six specific failure modes in the current landscape:

2.1 Data Extraction Without Compensation


When community members complete mental health assessments, financial literacy modules, or behavioral surveys through any conventional platform, the resulting data is treated as the platform's property. Academic researchers publish findings. Companies build products. Members receive nothing. This is not a minor ethical oversight — it is a structural feature of the current system that CIN is designed to dismantle.

2.2 Siloing of Community Knowledge

Behavioral and financial data collected from working-class and marginalized communities is especially valuable precisely because it is underrepresented in mainstream research. But because it is collected in silos — by individual apps, academic studies, and nonprofit programs — it cannot contribute to the kind of large-scale, longitudinal research that produces durable scientific insight. Community knowledge is fragmented by design, keeping it less actionable than centralized datasets controlled by major institutions [2].

2.3 Misaligned Incentive Structures

Current research incentives reward publication and data accumulation, not participant welfare or community benefit. Members who participate in studies receive, at best, modest one-time payments. They bear the costs of participation — time, vulnerability, disclosure — while institutions capture the long-term value. This misalignment actively discourages authentic engagement and produces data of lower quality than a genuinely incentive-aligned system would generate [3].



2.4 Lack of Participant Agency


Research participants in conventional studies have no ongoing relationship with the data they generate. They cannot see how it is used, cannot withdraw it from future analyses, cannot update it as their circumstances change, and cannot benefit from it as its value compounds over time. CIN treats data ownership as a continuous right, not a one-time transaction.

2.5 Concentration of AI Development

The behavioral AI systems that inform mental health apps, financial tools, and wellness platforms are developed on data sourced from communities but optimized for institutional goals. The models that result encode institutional biases, lack cultural specificity, and are unavailable for community-directed use. CIN enables MWA's community to contribute to AI model development in ways that keep both the process and the outputs aligned with community interests [4].

2.6 Privacy Vulnerability

Mental health and financial data are among the most sensitive categories of personal information. Centralized collection creates both technical vulnerabilities — data breach risk — and structural ones, where access to sensitive data by institutions enables discrimination, predatory targeting, and misuse. Decentralized data stores and federated learning eliminate the single-point-of-failure model that makes centralized collection so dangerous.



3. Background & Related Work

CIN builds on a growing body of work at the intersection of federated learning, blockchain coordination, and decentralized data ownership. This section situates CIN's contributions within that literature and identifies the specific gaps it addresses.

3.1 Federated Learning


Federated learning (FL) enables AI model training across distributed data sources without requiring data to be centralized [5]. Participating nodes — in CIN's case, member devices — train on local data and share only model parameter updates, not raw data. A global model is iteratively improved by aggregating these updates. Bonawitz et al. (2019) demonstrated that federated learning architectures can scale to hundreds of millions of devices while maintaining strong privacy guarantees [5]. CIN adapts this approach specifically for the behavioral and financial data generated through MWA's curriculum and assessment infrastructure.

Existing FL implementations, however, largely operate within institutional frameworks where a central aggregator — the research institution or platform — retains coordination authority. This reintroduces centralization at the protocol layer even when data remains distributed. CIN addresses this by moving coordination to a public blockchain smart contract, eliminating any single coordinating authority.

3.2 Blockchain-Coordinated FL

Several prior works have explored replacing central FL servers with blockchain-based coordination. Blockchain Assisted Federated Learning (BC-FL) by Wang et al. (2020) uses a public blockchain to coordinate model updates [6]. BlockFlow by Mugunthan et al. (2020) introduces an accountable, privacy-preserving FL protocol with on-chain reward distribution [7]. 2CP by Cai et al. (2020) proposes decentralized protocols for transparently evaluating contributor performance [8].

CIN extends these frameworks in two ways. First, it applies them to the specific context of behavioral health and financial research data — a domain with distinct privacy requirements and community ownership considerations. Second, it integrates the federated learning and rewards architecture directly with a structured educational platform (Blue and MWA's curriculum),




creating a seamless member experience in which participation in research is a natural extension of learning.

3.3 Decentralized Data Stores

Personal Data Stores (PDS) and related technologies [9][10] enable individuals to maintain sovereign control over their data across applications. Emerging protocols including Solid (developed by Tim Berners-Lee's team at MIT) and decentralized identity standards (DIDs, W3C) provide technical infrastructure for data sovereignty at the individual level. CIN's architecture is compatible with these emerging standards, allowing member data to remain in member-controlled stores while still contributing to federated model training.

3.4 Behavioral Health Research Infrastructure

Behavioral and mental health research faces particular challenges around participant trust, data sensitivity, and community representation. Research in this domain has documented systematic underrepresentation of working-class, BIPOC, and economically marginalized communities in datasets that are then used to build tools deployed in those same communities [11]. CIN is designed specifically to address this gap by enabling research that is community-directed, community-compensated, and community-owned.



4. MWA Platform Architecture

MWA's platform consists of four integrated layers that together enable the CIN protocol. Each layer is designed to be accessible to members with no technical background while maintaining the cryptographic guarantees necessary for trustless operation.

4.1 Curriculum & Assessment Layer


Members engage with MWA through structured learning tracks across three domains: mental wellness, financial literacy, and behavioral science. Curriculum is scaffolded progressively, with assessments built into the learning experience rather than appended as extractive add-ons. Validated instruments — including adapted versions of the PHQ-9, GAD-7, financial stress indices, and behavioral economics assessments — are embedded within modules in ways that generate research-grade data while providing members with personally meaningful feedback.

This layer is where member data is generated. The key architectural decision is that data generated here never leaves the member's device or personal data store without explicit, granular, and revocable consent. Members choose which data contributes to which research questions, and this choice is enforced at the protocol level rather than through policy alone.

4.2 Blue: AI Guide & Interface

Blue is MWA's AI guide — a conversational interface that accompanies members through curriculum, answers questions, supports reflection, and facilitates the kinds of structured behavioral conversations that generate high-quality research data. Blue is not simply a chatbot. It is the human-facing layer of the CIN infrastructure, designed to build the trust and longitudinal engagement that makes behavioral research meaningful.

Blue's interactions are structured to serve dual purposes: providing members with genuine learning and support, while generating analyzable behavioral data that contributes to the federated models. This dual purpose is transparent to members, who understand that their interactions with Blue contribute to community-owned research. Blue's conversational approach is informed by motivational interviewing principles, trauma-informed practice, and behavioral economics research on decision-making under financial stress.



4.3 Decentralized Data Store Layer


Each member's data is stored in a member-controlled decentralized data store — either on-device or in a personal data store of the member's choosing. Raw behavioral, assessment, and interaction data never leaves this store without the member's consent. Only model parameter updates — mathematical derivatives of local training — are shared with the network during federated learning rounds.

This architecture means that a breach of any single member's device does not compromise other members' data. It also means that MWA itself cannot access member data without permission — a structural privacy guarantee, not a policy one.

4.4 Blockchain Coordination Layer

A public blockchain smart contract coordinates the federated learning process, manages Validator assignments, records model update hashes, and distributes rewards. The use of a public blockchain — rather than a private or permissioned chain — is a deliberate design choice. It ensures that no single entity, including MWA itself, can unilaterally control the coordination process, block participation, or manipulate reward distribution.

Off-chain storage via IPFS (InterPlanetary File System) handles model weights and encrypted test datasets, with content identifiers (CIDs) recorded on-chain to ensure integrity. This combination reduces on-chain transaction costs while preserving the verifiability guarantees that make the protocol trustless.



5. Community Intelligence Network Protocol

The CIN protocol operationalizes the architecture described in Section 4. It consists of three coordinated subsystems: federated learning orchestration, decentralized aggregation, and trustless reward distribution.

5.1 Key Participants


The protocol involves three distinct roles, each with defined responsibilities and incentive structures:

- **Members (Participants):** Community members who engage with MWA curriculum, complete assessments, and interact with Blue. Members own their data, consent to specific research contributions, and receive token rewards for participation. Members do not need to understand the underlying cryptography — their participation is mediated through Blue and MWA's interface.
- **Research Owners (Model Owners):** Researchers, MWA staff, or community-designated research committees who define research questions, deploy federated learning protocols via smart contract, and provide the reward pool that compensates participating members. Research Owners gain access to trained models and aggregate insights without ever accessing individual member data.
- **Validators:** Staked nodes responsible for aggregating member model updates, evaluating contribution quality, and confirming reward calculations. Validators are economically incentivized to behave honestly through a proof-of-stake mechanism with slashing penalties for misbehavior. In early CIN deployments, MWA may operate initial Validator nodes, but the architecture is designed to support third-party and community Validators as the network matures.

5.2 Training Round Protocol

A complete CIN training round proceeds as follows:

1. Research Owner deploys an intelligence smart contract on the public blockchain, specifying the research parameters, the base model (stored on IPFS), and the reward pool deposited in the contract.




2. Members who have opted into the relevant research topic receive notification through Blue. Blue explains the research purpose, what data will be used, and what the member will receive in return — in plain language, not legalese.
3. Consenting members download the base model to their device. Using their locally stored behavioral and assessment data, their device runs training iterations, generating updated model parameters without exposing raw data.
4. Members upload encrypted model updates to IPFS, recording content identifiers on the smart contract. Raw data remains on-device.
5. Validators, randomly assigned to member subgroups, retrieve encrypted model updates and aggregate them using their own compute resources. Multiple Validators per subgroup ensure that no single Validator can corrupt the aggregation without detection.
6. The aggregated global model is published to IPFS. The smart contract verifies that Validator results are consistent. Discrepancies trigger on-chain dispute resolution.
7. The Research Owner evaluates the updated global model against a held-out encrypted test dataset. Training continues until the model meets predefined performance criteria or the Research Owner's reward pool is exhausted.
8. In the final round, Validators evaluate each member's contribution using the full test dataset and contributivity scoring methods (Shapley Values, median scoring, or Research Owner-specified alternatives). Rewards are calculated and distributed by the smart contract without Research Owner involvement.

5.3 Decentralized Auditing & Reward Distribution

The reward mechanism is designed around a single principle: members should receive fair compensation for their contributions without having to trust any institution to calculate or deliver it. Smart contracts calculate reward fractions based on verified contributivity scores. The Research Owner's deposited rewards are distributed automatically upon completion of the evaluation round.

Validators who submit evaluations that deviate significantly from the consensus — indicating dishonest evaluation — are penalized via token slashing. The smart contract enforces this automatically. Validators have strong economic incentives to evaluate honestly.


The test dataset used for evaluation is encrypted by the Research Owner and shared with Validators only after training is complete, preventing any participant from gaming the training process by optimizing for the evaluation set. A resampling strategy (drawing from a reserved 40% of available data) further reduces overfitting risk.



5.4 Blue as Research Interface

Blue's role in the CIN protocol goes beyond curriculum delivery. Blue is the member-facing interface through which consent is managed, research participation is explained, contribution status is tracked, and reward receipt is confirmed. This design decision is critical: the technical complexity of the underlying protocol should be invisible to members. Blue translates protocol events — a training round opening, a reward being distributed, a new research question becoming available — into plain-language interactions that members can act on without technical knowledge.

Blue also plays a data quality role. By building longitudinal relationships with members — tracking progress through curriculum, noting changes in assessment scores over time, and providing continuity across sessions — Blue enables the kind of rich, longitudinal behavioral data that is most valuable for research. Members who trust Blue and engage consistently produce better data than members recruited for one-time participation in a decontextualized study.



6. Research Applications

CIN is designed to support a specific class of research questions: those concerning the relationships between mental wellness, financial behavior, and community-level resilience, studied longitudinally in communities that are typically underrepresented in academic research. The following use cases illustrate the framework's scope:

6.1 Financial Stress & Mental Health Outcomes


MWA's curriculum integrates financial literacy and mental wellness in a way that reflects their real-world interdependence. CIN enables longitudinal research on how financial education interventions affect psychological wellbeing measures over time — a research question that requires exactly the kind of sustained, community-embedded data collection that conventional research cannot achieve. Members contribute assessment data across multiple curriculum modules, and the federated model learns associations between financial knowledge acquisition and changes in stress, anxiety, and self-efficacy measures without any individual's data leaving their device.

6.2 Behavioral Economics of Financial Decision-Making

Blue's structured conversations about financial decisions — budgeting choices, responses to financial emergencies, attitudes toward debt and savings — generate rich behavioral data about how real people under real financial pressure make financial decisions. This data, aggregated across the community, can produce insights about behavioral economic phenomena (loss aversion, hyperbolic discounting, social comparison effects) that are far more ecologically valid than laboratory studies, and more representative than datasets drawn from financially stable populations.


6.3 Community Resilience & Social Support Networks

MWA's community structure creates natural opportunities to study how social connection and peer support mediate the relationship between financial stress and psychological outcomes. CIN enables research on these network effects without requiring members to disclose sensitive information about their social relationships to any central authority.



6.4 AI Model Development for Underserved Communities

Perhaps most significantly, CIN enables the development of AI models — mental health screening tools, natural language interfaces for financial coaching, and behavioral prediction frameworks — that are natively calibrated to the structural contexts of underserved populations. Standard commercial models are structurally prone to demographic blindspots. By shifting model development to a community-owned federated loop, MWA can directly engineer and train localized AI architectures that maintain strict contextual relevance without sacrificing participant privacy.



7. Conclusion

The Community Intelligence Network represents a foundational shift away from the traditional, extraction-based paradigm of behavioral science and psychological data collection. By unifying decentralized data stores, blockchain orchestration, and federated learning, MWA creates an environment where data is treated as an active economic and intellectual asset owned entirely by the individual who produced it. Managed seamlessly through Blue, the underlying technical architecture ensures that community-driven research actively funds, represents, and empowers the community itself. Through CIN, the production of psychological and behavioral insights becomes what it always should have been: a collaborative, equitable, and completely transparent community enterprise.

